

Universidad de Oriente	
Machine Learning	Aplicación de algoritmos de clasificación para la detección temprana de abandono laboral

Descripción del Problema

El conjunto de datos de predicción de rotación de empleados contiene información de **10,000 empleados**, diseñado para el modelado predictivo y el análisis de la rotación laboral. Incluye una variedad de métricas demográficas, laborales y de desempeño que ayudan a comprender los factores que contribuyen a la salida o abandono de los empleados en una organización.

Variable objetivo: Attrition — Indica si el empleado dejó la empresa: Yes o No.

Se aplicarán los siguientes algoritmos de clasificación:

1. Regresión Logística
2. Algoritmo KNN (K-Nearest Neighbors)
3. Árbol de Decisiones

Para cada algoritmo se evaluarán **varios ejemplos** (variando hiperparámetros) a fin de definir la mejor respuesta predictiva.

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

from sklearn.model_selection import train_test_split
from sklearn.preprocessing import LabelEncoder, StandardScaler
from sklearn.metrics import accuracy_score, confusion_matrix,
    precision_score, recall_score, f1_score
from sklearn.linear_model import LogisticRegression
from sklearn.neighbors import KNeighborsClassifier
from sklearn.tree import DecisionTreeClassifier
from sklearn import tree
import dataframe_image as dfi

import warnings
warnings.filterwarnings('ignore')
```

1. Carga y Exploración del Conjunto de Datos

```
# Cargar el conjunto de datos
df = pd.read_csv('clasificacion.csv')
df.head()
```

	Employee_ID	Age	Gender	Marital_Status	Department	Job_Role	Jo
0	1	58	Female	Married	IT	Manager	1
1	2	48	Female	Married	Sales	Assistant	5
2	3	34	Male	Married	Marketing	Assistant	1
3	4	27	Female	Divorced	Marketing	Manager	1
4	5	40	Male	Divorced	Marketing	Executive	1

5 rows × 26 columns

Dimensiones del dataset

```
print(f"Filas: {df.shape[0]}, Columnas: {df.shape[1]}")
```

Filas: 1000, Columnas: 26

Información general del dataset

```
df.info()
```

```
<class 'pandas.DataFrame'>
```

```
RangeIndex: 1000 entries, 0 to 999
```

```
Data columns (total 26 columns):
```

```
#   Column                                Non-Null Count  Dtype
---  -
0   Employee_ID                          1000 non-null  int64
1   Age                                  1000 non-null  int64
2   Gender                               1000 non-null  str
3   Marital_Status                       1000 non-null  str
4   Department                           1000 non-null  str
5   Job_Role                             1000 non-null  str
6   Job_Level                            1000 non-null  int64
7   Monthly_Income                       1000 non-null  int64
8   Hourly_Rate                          1000 non-null  int64
9   Years_at_Company                     1000 non-null  int64
10  Years_in_Current_Role                 1000 non-null  int64
11  Years_Since_Last_Promotion             1000 non-null  int64
12  Work_Life_Balance                     1000 non-null  int64
13  Job_Satisfaction                       1000 non-null  int64
14  Performance_Rating                    1000 non-null  int64
15  Training_Hours_Last_Year              1000 non-null  int64
16  Overtime                              1000 non-null  str
17  Project_Count                         1000 non-null  int64
18  Average_Hours_Worked_Per_Week         1000 non-null  int64
19  Absenteeism                           1000 non-null  int64
20  Work_Environment_Satisfaction          1000 non-null  int64
21  Relationship_with_Manager              1000 non-null  int64
22  Job_Involvement                       1000 non-null  int64
23  Distance_From_Home                    1000 non-null  int64
24  Number_of_Companies_Worked             1000 non-null  int64
25  Attrition                             1000 non-null  str
```

```
dtypes: int64(20), str(6)
```

```
memory usage: 203.3 KB
```

```
# Estadísticas descriptivas de las variables numéricas
```

```
df.describe()
```

	Employee_ID	Age	Job_Level	Monthly_Income	Hourly_Rate
count	1000.000000	1000.000000	1000.000000	1000.000000	1000.000000
mean	500.500000	39.991000	3.055000	11499.899000	57.837000
std	288.819436	11.780055	1.399977	4920.529231	24.702031
min	1.000000	20.000000	1.000000	3001.000000	15.000000
25%	250.750000	30.000000	2.000000	7395.750000	36.000000
50%	500.500000	41.000000	3.000000	11256.000000	58.000000
75%	750.250000	50.250000	4.000000	15855.000000	80.000000
max	1000.000000	59.000000	5.000000	19999.000000	99.000000

```
# Distribución de la variable objetivo Attrition
```

```
print(df['Attrition'].value_counts())
```

```
print()
```

```
print(df['Attrition'].value_counts(normalize=True) * 100)
```

```
Attrition
```

```
No      811
```

```
Yes      189
```

```
Name: count, dtype: int64
```

```
Attrition
```

```
No      81.1
```

```
Yes      18.9
```

```
Name: proportion, dtype: float64
```

```
# Gráfico de distribución de la variable objetivo
```

```
plt.figure(figsize=(6, 4))
```

```
ax = sns.countplot(x='Attrition', data=df)
```

```
total = len(df)
```

```
for p in ax.patches:
```

```
    cantidad = int(p.get_height())
```

```
    porcentaje = cantidad / total * 100
```


```
    ax.annotate(f'{cantidad}\n{porcentaje:.1f}%',  
               (p.get_x() + p.get_width() / 2, p.get_height()),  
               ha='center', va='bottom')
```

```
plt.title('Distribución de la variable Attrition', y=1.05)
```

```
plt.xlabel('Attrition')
```

```
plt.ylabel('Cantidad de empleados')
```

```
plt.show()
```




2. Análisis Exploratorio de Datos (EDA)

Histogramas de las variables numéricas

```
numeric_cols = df.select_dtypes(include=[np.number]).columns.tolist()
plt.figure()
df[numeric_cols].hist(figsize=(14, 10), edgecolor='black')
plt.suptitle('Distribución de variables numéricas', y=1.02)
plt.tight_layout()
plt.show()
```

<Figure size 640x480 with 0 Axes>

Distribución de variables numéricas




Boxplots de variables clave según Attrition

```
fig, axes = plt.subplots(2, 3, figsize=(14, 8))
```

```
vars_to_plot = ['Age', 'Monthly_Income', 'Years_at_Company',
                'Job_Satisfaction', 'Work_Life_Balance',
                'Distance_From_Home']
```

```
for ax, var in zip(axes.flatten(), vars_to_plot):
    sns.boxplot(x='Attrition', y=var, data=df, ax=ax)
    ax.set_title(f'{var} por Attrition')
```

```
plt.tight_layout()
plt.show()
```



Attrition por variables categóricas


```
fig, axes = plt.subplots(1, 3, figsize=(16, 5))
```

```
for ax, col in zip(axes, ['Department', 'Gender', 'Overtime']):
    attrition_rate = df.groupby(col)['Attrition'].apply(
        lambda x: (x == 'Yes').mean() * 100
    ).reset_index()
    attrition_rate.columns = [col, 'Attrition_Rate']
    sns.barplot(x=col, y='Attrition_Rate', data=attrition_rate, ax=ax)
    for p in ax.patches:
        porcentaje = p.get_height()
        ax.annotate(f'{porcentaje:.1f}%',
                    (p.get_x() + p.get_width() / 2, p.get_height()),
                    ha='center', va='bottom')
    ax.set_title(f'Tasa de Attrition por {col} (%)')
    ax.set_ylabel('Tasa de abandono (%)')
    ax.tick_params(axis='x', rotation=15)
```

```
fig.suptitle('Tasa de abandono por variables categóricas')
```

```
plt.tight_layout()
```

```
plt.show()
```



3. Preprocesamiento de Datos

Eliminar columna identificadora (no aporta información predictiva)

```
df_model = df.drop(columns=['Employee_ID'])
```

Codificar variables categóricas con LabelEncoder

```
le = LabelEncoder()
```

```
categorical_cols = df_model.select_dtypes(include=[  
    'object']).columns.tolist()
```

```
print("Variables categóricas a codificar:", categorical_cols)
```

```
for col in categorical_cols:
```

```
    df_model[col] = le.fit_transform(df_model[col])
```

```
df_model.head()
```

Variables categóricas a codificar: ['Gender', 'Marital_Status', 'Department', 'Job_Role', 'Overtime', 'Attrition']

	Age	Gender	Marital_Status	Department	Job_Role	Job_Level	Monthly_Income
0	58	0	1	2	3	1	15488
1	48	0	1	4	1	5	13079
2	34	1	1	3	1	1	13744
3	27	0	0	3	3	1	6809
4	40	1	0	3	2	1	10206

5 rows × 25 columns

Separar variables predictoras (X) y variable objetivo (y)

```
X = df_model.drop(columns=['Attrition'])
```

```
y = df_model['Attrition']
```

Dividir en conjunto de entrenamiento (80%) y prueba (20%)

```
X_train, X_test, y_train, y_test = train_test_split(  
    X, y, test_size=0.20, random_state=2026
```

```
)
```

```
print(f"Tamaño entrenamiento: {X_train.shape}")
```

```
print(f"Tamaño prueba: {X_test.shape}")
```

Tamaño entrenamiento: (800, 24)

Tamaño prueba: (200, 24)

Escalar las variables (necesario para Regresión Logística y KNN)

```
scaler = StandardScaler()
```

```
X_train_scaled = scaler.fit_transform(X_train)
```

```
X_test_scaled = scaler.transform(X_test)
```

4. Regresión Logística


La **Regresión Logística** es un modelo de clasificación que estima la probabilidad de que una observación pertenezca a una clase determinada. Utiliza la función sigmoide para mapear valores reales al intervalo $[0, 1]$.

Se evaluarán tres configuraciones distintas:

- a) Parámetros por defecto
- b) Regularización L1 (lasso)
- c) Mayor número de iteraciones con penalización L2

a) Regresión Logística con parámetros por defecto

```
# Crear el objeto de clase Regresión Logística con parámetros por defecto
lr_a = LogisticRegression(random_state=2026)
# Ajustar el modelo
lr_a.fit(X_train_scaled, y_train)
```



```
# Obtener predicciones
y_pred_lr_a = lr_a.predict(X_test_scaled)

# Accuracy Score
accuracy_lr_a = accuracy_score(y_test, y_pred_lr_a)
print("El accuracy score que se tiene es %s" % accuracy_lr_a)
```


El accuracy score que se tiene es 0.785

```
from sklearn.metrics import confusion_matrix
```

```
# Obtener la matriz de confusión
cm = confusion_matrix(y_test, y_pred_lr_a)
sns.heatmap(cm, annot=True, fmt="d")

plt.title("Regresión Logística - Defecto")
plt.xlabel("Predicho")
plt.ylabel("Real")
```

```
Text(50.72222222222214, 0.5, 'Real')
```



Obtener la precisión, recall y f1

```
precision_lr_a = precision_score(y_test, y_pred_lr_a, average="macro")
```

```
recall_lr_a = recall_score(y_test, y_pred_lr_a, average="macro")
```

```
f1_lr_a = f1_score(y_test, y_pred_lr_a, average="macro")
```

```
metricas = pd.DataFrame({"Métricas": ["Precisión", "Recall", "F1"],
                          "Valores": [precision_lr_a, recall_lr_a, f1_lr_a]})
```

metricas

	Métricas	Valores
0	Precisión	0.392500
1	Recall	0.500000
2	F1	0.439776

b) Regresión Logística con regularización L1 (Lasso)

La regularización L1 produce modelos más dispersos al forzar algunos coeficientes a cero, actuando como un selector automático de variables.

Crear el objeto de clase Regresión Logística con regularización L1

```
lr_b = LogisticRegression(penalty='l1', solver='liblinear',
                          random_state=2026)
```

Ajustar el modelo

```
lr_b.fit(X_train_scaled, y_train)
```

▼ LogisticRegression ⓘ ?

► Parameters

Obtener predicciones

```
y_pred_lr_b = lr_b.predict(X_test_scaled)
```

Accuracy Score

```
accuracy_lr_b = accuracy_score(y_test, y_pred_lr_b)
```

```
print("El accuracy score que se tiene es %s" % accuracy_lr_b)
```

El accuracy score que se tiene es 0.785

Obtener La matriz de confusión

```
cm = confusion_matrix(y_test, y_pred_lr_b)
```


```
sns.heatmap(cm, annot=True, fmt="d")
```

```
plt.title("Regresión Logística - L1 (Lasso)")
```

```
plt.xlabel("Predicho")
```

```
plt.ylabel("Real")
```

```
Text(50.72222222222214, 0.5, 'Real')
```



Obtener la precisión, recall y f1

```
precision_lr_b = precision_score(y_test, y_pred_lr_b, average="macro")
```

```
recall_lr_b = recall_score(y_test, y_pred_lr_b, average="macro")
```

```
f1_lr_b = f1_score(y_test, y_pred_lr_b, average="macro")
```

```
metricas = pd.DataFrame({"Métricas": ["Precisión", "Recall", "F1"],
                        "Valores": [precision_lr_b, recall_lr_b, f1_lr_b]})

metricas


#dfi.export(metricas, "metricas_arbol_max_depth_5.png")
```

	Métricas	Valores	
0	Precisión	0.392500	
1	Recall	0.500000	
2	F1	0.439776	

c) Regresión Logística con mayor número de iteraciones y C=0.5

Reducir el valor de C aumenta la regularización, lo que puede mejorar la generalización del modelo.

```
# Crear el objeto de clase Regresión Logística con C=0.5 y max_iter=500
lr_c = LogisticRegression(C=0.5, max_iter=500, random_state=2026)
# Ajustar el modelo
lr_c.fit(X_train_scaled, y_train)
```




```
# Obtener predicciones
y_pred_lr_c = lr_c.predict(X_test_scaled)

# Accuracy Score
accuracy_lr_c = accuracy_score(y_test, y_pred_lr_c)
print("El accuracy score que se tiene es %s" % accuracy_lr_c)
```

El accuracy score que se tiene es 0.785

```
# Obtener la matriz de confusión
cm = confusion_matrix(y_test, y_pred_lr_c)
sns.heatmap(cm, annot=True, fmt="d")
plt.title("Regresión Logística - C=0.5")
plt.xlabel("Predicho")
plt.ylabel("Real")

Text(50.72222222222214, 0.5, 'Real')
```



Obtener la precisión, recall y f1

```
precision_lr_c = precision_score(y_test, y_pred_lr_c, average="macro")
recall_lr_c = recall_score(y_test, y_pred_lr_c, average="macro")
f1_lr_c = f1_score(y_test, y_pred_lr_c, average="macro")

metricas = pd.DataFrame({"Métricas": ["Precisión", "Recall", "F1"],
                          "Valores": [precision_lr_c, recall_lr_c, f1_lr_c]})

metricas
```

	Métricas	Valores
0	Precisión	0.392500
1	Recall	0.500000
2	F1	0.439776

Comparación de los modelos de Regresión Logística

A continuación se resumen las métricas de los tres modelos de Regresión Logística evaluados:

```
comparacion_lr = pd.DataFrame({
    'Modelo': ['RL - Defecto', 'RL - L1 (Lasso)', 'RL - C=0.5, iter=500'],
    'Accuracy': [accuracy_lr_a, accuracy_lr_b, accuracy_lr_c],
    'Precisión': [precision_lr_a, precision_lr_b, precision_lr_c],
    'Recall': [recall_lr_a, recall_lr_b, recall_lr_c],
})
```

```

    'F1': [f1_lr_a, f1_lr_b, f1_lr_c]
})
comparacion_lr

```

	Modelo	Accuracy	Precisión	Recall	F1
0	RL - Defecto	0.785	0.3925	0.5	0.439776
1	RL - L1 (Lasso)	0.785	0.3925	0.5	0.439776
2	RL - C=0.5, iter=500	0.785	0.3925	0.5	0.439776

5. K-Nearest Neighbors (KNN)

El algoritmo **KNN** clasifica una observación en función de la clase mayoritaria entre sus k vecinos más cercanos en el espacio de características.

Se evaluarán tres configuraciones:


- a) k=5 (valor por defecto)
- b) k=3
- c) k=11

a) KNN con k=5 (parámetros por defecto)

```

# Crear el objeto de clase KNN con parámetros por defecto (k=5)
knn_a = KNeighborsClassifier()
# Ajustar el modelo
knn_a.fit(X_train_scaled, y_train)

```



```

# Obtener el accuracy score
y_pred_knn_a = knn_a.predict(X_test_scaled)
accuracy_knn_a = accuracy_score(y_test, y_pred_knn_a)
print("El accuracy score que se tiene es %s" % accuracy_knn_a)

```


El accuracy score que se tiene es 0.77

```

# Obtener la matriz de confusión
cm = confusion_matrix(y_test, y_pred_knn_a)
sns.heatmap(cm, annot=True, fmt="d")
plt.title("KNN - k=5")
plt.xlabel("Predicho")
plt.ylabel("Real")

```

```
Text(50.72222222222214, 0.5, 'Real')
```



Obtener la precisión, recall y f1

```
precision_knn_a = precision_score(y_test, y_pred_knn_a, average="macro")
recall_knn_a = recall_score(y_test, y_pred_knn_a, average="macro")
f1_knn_a = f1_score(y_test, y_pred_knn_a, average="macro")
```

```
metricas = pd.DataFrame({"Métricas": ["Precisión", "Recall", "F1"],
                          "Valores": [precision_knn_a, recall_knn_a,
                                      f1_knn_a]})
```

metricas

	Métricas	Valores
0	Precisión	0.492308
1	Recall	0.498889
2	F1	0.455492

b) KNN con k=3

Con un valor de k más pequeño, el modelo es más sensible a las variaciones locales, lo que puede producir fronteras de decisión más irregulares.

Crear el objeto de clase KNN con k=3

```
knn_b = KNeighborsClassifier(n_neighbors=3)
```

Ajustar el modelo

```
knn_b.fit(X_train_scaled, y_train)
```

▼ KNeighborsClassifier ⓘ ?

▶ Parameters

Obtener el accuracy score


```
y_pred_knn_b = knn_b.predict(X_test_scaled)
accuracy_knn_b = accuracy_score(y_test, y_pred_knn_b)
print("El accuracy score que se tiene es %s" % accuracy_knn_b)
```

El accuracy score que se tiene es 0.74

Obtener la matriz de confusión

```
cm = confusion_matrix(y_test, y_pred_knn_b)
sns.heatmap(cm, annot=True, fmt="d")
plt.title("KNN - k=3")
plt.xlabel("Predicho")
plt.ylabel("Real")
```

```
Text(50.72222222222214, 0.5, 'Real')
```



Obtener la precisión, recall y f1

```
precision_knn_b = precision_score(y_test, y_pred_knn_b, average="macro")
recall_knn_b = recall_score(y_test, y_pred_knn_b, average="macro")
f1_knn_b = f1_score(y_test, y_pred_knn_b, average="macro")
```

```
metricas = pd.DataFrame({"Métricas": ["Precisión", "Recall", "F1"],
```

```

        "Valores": [precision_knn_b, recall_knn_b,
                    f1_knn_b]))
metricas

```

	Métricas	Valores
0	Precisión	0.491892
1	Recall	0.496667
2	F1	0.475701


c) KNN con k=11

Con un valor de k más grande, el modelo suaviza las fronteras de decisión y puede ser más robusto al ruido en los datos.

```

# Crear el objeto de clase KNN con k=11
knn_c = KNeighborsClassifier(n_neighbors=11)
# Ajustar el modelo
knn_c.fit(X_train_scaled, y_train)

```



```

# Obtener el accuracy score
y_pred_knn_c = knn_c.predict(X_test_scaled)
accuracy_knn_c = accuracy_score(y_test, y_pred_knn_c)
print("El accuracy score que se tiene es %s" % accuracy_knn_c)

```

El accuracy score que se tiene es 0.78

```


# Obtener la matriz de confusión
cm = confusion_matrix(y_test, y_pred_knn_c)
sns.heatmap(cm, annot=True, fmt="d")
plt.title("KNN - k=11")
plt.xlabel("Predicho")
plt.ylabel("Real")

```

```

Text(50.72222222222214, 0.5, 'Real')

```



Obtener la precisión, recall y f1

```
precision_knn_c = precision_score(y_test, y_pred_knn_c, average="macro")
recall_knn_c = recall_score(y_test, y_pred_knn_c, average="macro")
f1_knn_c = f1_score(y_test, y_pred_knn_c, average="macro")
```

```
metricas = pd.DataFrame({"Métricas": ["Precisión", "Recall", "F1"],
                          "Valores": [precision_knn_c, recall_knn_c,
                                      f1_knn_c]})
```

metricas

	Métricas	Valores
0	Precisión	0.391960
1	Recall	0.496815
2	F1	0.438202

Comparación de los modelos KNN

```
comparacion_knn = pd.DataFrame({
    'Modelo': ['KNN - k=5 (defecto)', 'KNN - k=3', 'KNN - k=11'],
    'Accuracy': [accuracy_knn_a, accuracy_knn_b, accuracy_knn_c],
    'Precisión': [precision_knn_a, precision_knn_b, precision_knn_c],
    'Recall': [recall_knn_a, recall_knn_b, recall_knn_c],
    'F1': [f1_knn_a, f1_knn_b, f1_knn_c]
})
```

```
})
comparacion_knn
```

	Modelo	Accuracy	Precisión	Recall	F1
0	KNN - k=5 (defecto)	0.77	0.492308	0.498889	0.455492
1	KNN - k=3	0.74	0.491892	0.496667	0.475701
2	KNN - k=11	0.78	0.391960	0.496815	0.438202

6.Árbol de Decisiones


Un **árbol de decisión** consiste en una secuencia de condiciones sobre las distintas variables predictoras y su relación con la variable de respuesta. Estas condiciones se van sucediendo en distintos caminos del árbol, dividiendo el espacio de los predictores en diferentes regiones.

Se evaluarán tres configuraciones:

- a) Parámetros por defecto
- b) Profundidad máxima max_depth=5
- c) Mínimo de muestras en hoja min_samples_leaf=4

a) Árbol de Decisiones con parámetros por defecto

```
# Crear el objeto de clase árbol
tree_clf_a = DecisionTreeClassifier(random_state=2026)
# Ajustar el árbol
tree_clf_a.fit(X_train, y_train)
```




```
# Obtener el accuracy score
y_pred_a = tree_clf_a.predict(X_test)
accuracy_a = accuracy_score(y_test, y_pred_a)
print("El accuracy score que se tiene es %s" % accuracy_a)
```

El accuracy score que se tiene es 0.68

```
# Obtener la matriz de confusión
cm = confusion_matrix(y_test, y_pred_a)
sns.heatmap(cm, annot=True, fmt="d")
plt.title("Árbol de Decisiones - Defecto")
plt.xlabel("Predicho")
plt.ylabel("Real")
```

```
Text(50.72222222222214, 0.5, 'Real')
```



Obtener la precisión, recall y f1

```
precision_a = precision_score(y_test, y_pred_a, average="macro")
```

```
recall_a = recall_score(y_test, y_pred_a, average="macro")
```

```
f1_a = f1_score(y_test, y_pred_a, average="macro")
```

```
metricas = pd.DataFrame({"Métricas": ["Precisión", "Recall", "F1"],
                          "Valores": [precision_a, recall_a, f1_a]})
```

```
metricas
```

	Métricas	Valores
0	Precisión	0.564890
1	Recall	0.576655
2	F1	0.567509

b) Árbol de Decisiones con profundidad máxima `max_depth=5`

Una manera sencilla de evitar el posible sobreajuste del árbol es fijar un número de profundidad máxima, es decir, del número de divisiones sucesivas que se realizan.

Crear el objeto de clase árbol con `max_depth=5`

```
tree_clf_b = DecisionTreeClassifier(max_depth=5, random_state=2026)
```

Ajustar el árbol

```
tree_clf_b.fit(X_train, y_train)
```

```
plt.figure(figsize=(14,10))
```

<Figure size 1400x1000 with 0 Axes>

<Figure size 1400x1000 with 0 Axes>

Visualizar el árbol de decisión

tree.plot_tree(tree_clf_b, filled=True)


```
[Text(0.5703125, 0.9166666666666666, 'x[10] <= 6.5\ngini = 0.298\nsamples = 800\nvalue = [654, 146]'),
 Text(0.2864583333333333, 0.75, 'x[14] <= 56.5\ngini = 0.325\nsamples = 553\nvalue = [440, 113]'),
 Text(0.42838541666666663, 0.8333333333333333, 'True '),
 Text(0.15625, 0.5833333333333334, 'x[5] <= 3.5\ngini = 0.36\nsamples = 306\nvalue = [234.0, 72.0]'),
 Text(0.08333333333333333, 0.4166666666666667, 'x[18] <= 4.5\ngini = 0.408\nsamples = 175\nvalue = [125, 50]'),
 Text(0.041666666666666664, 0.25, 'x[12] <= 2.5\ngini = 0.236\nsamples = 44\nvalue = [38, 6]'),
 Text(0.020833333333333332, 0.08333333333333333, 'gini = 0.473\nsamples = 13\nvalue = [8, 5]'),
 Text(0.0625, 0.08333333333333333, 'gini = 0.062\nsamples = 31\nvalue = [30, 1]'),
 Text(0.125, 0.25, 'x[20] <= 2.5\ngini = 0.446\nsamples = 131\nvalue = [87, 44]'),
 Text(0.10416666666666667, 0.08333333333333333, 'gini = 0.375\nsamples = 68\nvalue = [51, 17]'),
 Text(0.14583333333333334, 0.08333333333333333, 'gini = 0.49\nsamples = 63\nvalue = [36, 27]'),
 Text(0.22916666666666666, 0.4166666666666667, 'x[6] <= 19878.5\ngini = 0.279\nsamples = 131\nvalue = [109, 22]'),
 Text(0.20833333333333334, 0.25, 'x[14] <= 52.5\ngini = 0.262\nsamples = 129\nvalue = [109, 20]'),
 Text(0.1875, 0.08333333333333333, 'gini = 0.225\nsamples = 116\nvalue = [101.0, 15.0]'),
 Text(0.22916666666666666, 0.08333333333333333, 'gini = 0.473\nsamples = 13\nvalue = [8, 5]'),
 Text(0.25, 0.25, 'gini = 0.0\nsamples = 2\nvalue = [0, 2]'),
 Text(0.4166666666666667, 0.5833333333333334, 'x[8] <= 6.5\ngini = 0.277\nsamples = 247\nvalue = [206, 41]'),
 Text(0.3333333333333333, 0.4166666666666667, 'x[8] <= 2.5\ngini = 0.44\nsamples = 52\nvalue = [35, 17]'),
 Text(0.2916666666666667, 0.25, 'x[10] <= 4.5\ngini = 0.198\nsamples = 18\nvalue = [16, 2]'),
 Text(0.27083333333333333, 0.08333333333333333, 'gini = 0.0\nsamples = 13\nvalue = [13, 0]'),
 Text(0.3125, 0.08333333333333333, 'gini = 0.48\nsamples = 5\nvalue = [3, 2]'),
 Text(0.375, 0.25, 'x[9] <= 6.5\ngini = 0.493\nsamples = 34\nvalue = [19, 15]'),
 Text(0.3541666666666667, 0.08333333333333333, 'gini = 0.337\nsamples = 14\nvalue = [11, 3]'),
 Text(0.39583333333333333, 0.08333333333333333, 'gini = 0.48\nsamples = 20\nvalue = [8, 12]'),
 Text(0.5, 0.4166666666666667, 'x[7] <= 26.5\ngini = 0.216\nsamples =
```

```
195\nvalue = [171, 24]'),
  Text(0.4583333333333333, 0.25, 'x[14] <= 91.0\ngini = 0.375\nsamples =
32\nvalue = [24, 8]'),
  Text(0.4375, 0.0833333333333333, 'gini = 0.285\nsamples = 29\nvalue = [24,
5]'),
  Text(0.4791666666666667, 0.0833333333333333, 'gini = 0.0\nsamples =
3\nvalue = [0, 3]'),
  Text(0.5416666666666666, 0.25, 'x[5] <= 3.5\ngini = 0.177\nsamples =
163\nvalue = [147, 16]'),
  Text(0.5208333333333334, 0.0833333333333333, 'gini = 0.095\nsamples =
100\nvalue = [95, 5]'),
  Text(0.5625, 0.0833333333333333, 'gini = 0.288\nsamples = 63\nvalue = [52,
11]'),
  Text(0.8541666666666666, 0.75, 'x[18] <= 18.5\ngini = 0.232\nsamples =
247\nvalue = [214, 33]'),
  Text(0.7122395833333333, 0.8333333333333333, ' False'),
  Text(0.75, 0.5833333333333334, 'x[22] <= 21.5\ngini = 0.207\nsamples =
230\nvalue = [203, 27]'),
  Text(0.6666666666666666, 0.4166666666666667, 'x[12] <= 3.5\ngini =
0.299\nsamples = 104\nvalue = [85, 19]'),
  Text(0.625, 0.25, 'x[16] <= 7.5\ngini = 0.192\nsamples = 65\nvalue = [58,
7]'),
  Text(0.6041666666666666, 0.0833333333333333, 'gini = 0.078\nsamples =
49\nvalue = [47.0, 2.0]'),
  Text(0.6458333333333334, 0.0833333333333333, 'gini = 0.43\nsamples =
16\nvalue = [11, 5]'),
  Text(0.7083333333333334, 0.25, 'x[7] <= 17.5\ngini = 0.426\nsamples =
39\nvalue = [27, 12]'),
  Text(0.6875, 0.0833333333333333, 'gini = 0.0\nsamples = 2\nvalue = [0,
2]'),
  Text(0.7291666666666666, 0.0833333333333333, 'gini = 0.394\nsamples =
37\nvalue = [27, 10]'),
  Text(0.8333333333333334, 0.4166666666666667, 'x[17] <= 31.5\ngini =
0.119\nsamples = 126\nvalue = [118, 8]'),
  Text(0.7916666666666666, 0.25, 'x[13] <= 1.5\ngini = 0.469\nsamples =
8\nvalue = [5, 3]'),
  Text(0.7708333333333334, 0.0833333333333333, 'gini = 0.0\nsamples =
2\nvalue = [0, 2]'),
  Text(0.8125, 0.0833333333333333, 'gini = 0.278\nsamples = 6\nvalue = [5,
1]'),
  Text(0.875, 0.25, 'x[10] <= 8.5\ngini = 0.081\nsamples = 118\nvalue = [113,
5]'),
  Text(0.8541666666666666, 0.0833333333333333, 'gini = 0.025\nsamples =
80\nvalue = [79, 1]'),
  Text(0.8958333333333334, 0.0833333333333333, 'gini = 0.188\nsamples =
38\nvalue = [34, 4]'),
  Text(0.9583333333333334, 0.5833333333333334, 'x[7] <= 80.5\ngini =
0.457\nsamples = 17\nvalue = [11, 6]'),
  Text(0.9375, 0.4166666666666667, 'x[11] <= 1.5\ngini = 0.26\nsamples =
13\nvalue = [11, 2]'),
  Text(0.9166666666666666, 0.25, 'gini = 0.0\nsamples = 1\nvalue = [0, 1]'),
  Text(0.9583333333333334, 0.25, 'x[19] <= 3.5\ngini = 0.153\nsamples =
12\nvalue = [11, 1]'),
  Text(0.9375, 0.0833333333333333, 'gini = 0.0\nsamples = 11\nvalue = [11,
```

```

0]'),
  Text(0.9791666666666666, 0.08333333333333333, 'gini = 0.0\nsamples =
1\nvalue = [0, 1]'),
  Text(0.9791666666666666, 0.4166666666666667, 'gini = 0.0\nsamples = 4\nvalue
= [0, 4]')]

```



```

# Obtener el accuracy score
y_pred_b = tree_clf_b.predict(X_test)
accuracy_b = accuracy_score(y_test, y_pred_b)
print("El accuracy score que se tiene es %s" % accuracy_b)

```

El accuracy score que se tiene es 0.765

```


# Obtener la matriz de confusión
cm = confusion_matrix(y_test, y_pred_b)
sns.heatmap(cm, annot=True, fmt="d")
plt.title("Árbol de Decisiones - max_depth=5")
plt.xlabel("Predicho")
plt.ylabel("Real")

```

```

Text(50.72222222222214, 0.5, 'Real')

```



Obtener la precisión, recall y f1

```
precision_b = precision_score(y_test, y_pred_b, average="macro")
```

```
recall_b = recall_score(y_test, y_pred_b, average="macro")
```

```
f1_b = f1_score(y_test, y_pred_b, average="macro")
```

```
metricas = pd.DataFrame({"Métricas": ["Precisión", "Recall", "F1"],
                          "Valores": [precision_b, recall_b, f1_b]})
```

```
metricas
```

	Métricas	Valores
0	Precisión	0.475086
1	Recall	0.495704
2	F1	0.453457

c) Árbol de Decisiones indicando que el mínimo número de observaciones que debe contener cada nodo hoja es de 4.


Otra manera de restringir la longitud del árbol de decisión sin fijar la profundidad máxima es indicando el número mínimo de observaciones que debe contener cada nodo hoja. En este caso se fijará el valor de `min_samples_leaf=4`.

Crear el objeto de clase árbol

```
tree_clf_c = DecisionTreeClassifier(min_samples_leaf=4, random_state=2026) #  
4 muestras
```

Ajustar el árbol

```
tree_clf_c.fit(X_train, y_train)
```



```
from sklearn import tree
```

```
tree.plot_tree(tree_clf_c, filled=True)
```

```
[Text(0.6492718446601942, 0.9583333333333334, 'x[10] <= 6.5\ngini =  
0.298\nsamples = 800\nvalue = [654, 146]'),  
Text(0.375, 0.875, 'x[14] <= 56.5\ngini = 0.325\nsamples = 553\nvalue =  
[440, 113]'),  
Text(0.5121359223300971, 0.9166666666666667, 'True '),  
Text(0.2275485436893204, 0.7916666666666666, 'x[5] <= 3.5\ngini =  
0.36\nsamples = 306\nvalue = [234.0, 72.0]'),  
Text(0.1104368932038835, 0.7083333333333334, 'x[18] <= 4.5\ngini =  
0.408\nsamples = 175\nvalue = [125, 50]'),  
Text(0.038834951456310676, 0.625, 'x[12] <= 2.5\ngini = 0.236\nsamples =  
44\nvalue = [38, 6]'),  
Text(0.019417475728155338, 0.5416666666666666, 'x[8] <= 15.5\ngini =  
0.473\nsamples = 13\nvalue = [8, 5]'),  
Text(0.009708737864077669, 0.4583333333333333, 'gini = 0.32\nsamples =  
5\nvalue = [1, 4]'),  
Text(0.02912621359223301, 0.4583333333333333, 'x[9] <= 6.5\ngini =  
0.219\nsamples = 8\nvalue = [7, 1]'),  
Text(0.019417475728155338, 0.375, 'gini = 0.375\nsamples = 4\nvalue = [3,  
1]'),  
Text(0.038834951456310676, 0.375, 'gini = 0.0\nsamples = 4\nvalue = [4,  
0]'),  
Text(0.05825242718446602, 0.5416666666666666, 'x[7] <= 28.0\ngini =  
0.062\nsamples = 31\nvalue = [30, 1]'),  
Text(0.04854368932038835, 0.4583333333333333, 'gini = 0.375\nsamples =  
4\nvalue = [3, 1]'),  
Text(0.06796116504854369, 0.4583333333333333, 'gini = 0.0\nsamples =  
27\nvalue = [27, 0]'),  
Text(0.1820388349514563, 0.625, 'x[20] <= 2.5\ngini = 0.446\nsamples =  
131\nvalue = [87, 44]'),  
Text(0.11165048543689321, 0.5416666666666666, 'x[8] <= 13.5\ngini =  
0.375\nsamples = 68\nvalue = [51, 17]'),  
Text(0.08737864077669903, 0.4583333333333333, 'x[8] <= 6.5\ngini =  
0.137\nsamples = 27\nvalue = [25, 2]'),  
Text(0.07766990291262135, 0.375, 'x[21] <= 2.5\ngini = 0.26\nsamples =  
13\nvalue = [11, 2]'),  
Text(0.06796116504854369, 0.2916666666666667, 'gini = 0.0\nsamples =  
8\nvalue = [8, 0]'),  
Text(0.08737864077669903, 0.2916666666666667, 'gini = 0.48\nsamples =  
5\nvalue = [3, 2]'),  
Text(0.0970873786407767, 0.375, 'gini = 0.0\nsamples = 14\nvalue = [14,
```

```
0]'),
  Text(0.13592233009708737, 0.4583333333333333, 'x[16] <= 5.5\ngini =
0.464\nsamples = 41\nvalue = [26, 15]'),
  Text(0.11650485436893204, 0.375, 'x[6] <= 7027.5\ngini = 0.49\nsamples =
21\nvalue = [9, 12]'),
  Text(0.10679611650485436, 0.2916666666666667, 'gini = 0.245\nsamples =
7\nvalue = [6, 1]'),
  Text(0.1262135922330097, 0.2916666666666667, 'x[14] <= 9.0\ngini =
0.337\nsamples = 14\nvalue = [3, 11]'),
  Text(0.11650485436893204, 0.2083333333333333, 'gini = 0.375\nsamples =
4\nvalue = [3, 1]'),
  Text(0.13592233009708737, 0.2083333333333333, 'gini = 0.0\nsamples =
10\nvalue = [0, 10]'),
  Text(0.1553398058252427, 0.375, 'x[6] <= 8283.0\ngini = 0.255\nsamples =
20\nvalue = [17, 3]'),
  Text(0.14563106796116504, 0.2916666666666667, 'gini = 0.49\nsamples =
7\nvalue = [4, 3]'),
  Text(0.1650485436893204, 0.2916666666666667, 'gini = 0.0\nsamples =
13\nvalue = [13, 0]'),
  Text(0.2524271844660194, 0.5416666666666666, 'x[3] <= 3.5\ngini =
0.49\nsamples = 63\nvalue = [36, 27]'),
  Text(0.24271844660194175, 0.4583333333333333, 'x[16] <= 5.5\ngini =
0.469\nsamples = 56\nvalue = [35, 21]'),
  Text(0.20388349514563106, 0.375, 'x[15] <= 0.5\ngini = 0.383\nsamples =
31\nvalue = [23, 8]'),
  Text(0.18446601941747573, 0.2916666666666667, 'x[8] <= 6.5\ngini =
0.5\nsamples = 14\nvalue = [7, 7]'),
  Text(0.17475728155339806, 0.2083333333333333, 'gini = 0.32\nsamples =
5\nvalue = [1, 4]'),
  Text(0.1941747572815534, 0.2083333333333333, 'x[6] <= 8181.5\ngini =
0.444\nsamples = 9\nvalue = [6, 3]'),
  Text(0.18446601941747573, 0.125, 'gini = 0.48\nsamples = 5\nvalue = [2,
3]'),
  Text(0.20388349514563106, 0.125, 'gini = 0.0\nsamples = 4\nvalue = [4, 0]'),
  Text(0.22330097087378642, 0.2916666666666667, 'x[12] <= 1.5\ngini =
0.111\nsamples = 17\nvalue = [16, 1]'),
  Text(0.21359223300970873, 0.2083333333333333, 'gini = 0.375\nsamples =
4\nvalue = [3, 1]'),
  Text(0.23300970873786409, 0.2083333333333333, 'gini = 0.0\nsamples =
13\nvalue = [13, 0]'),
  Text(0.2815533980582524, 0.375, 'x[19] <= 2.5\ngini = 0.499\nsamples =
25\nvalue = [12, 13]'),
  Text(0.2621359223300971, 0.2916666666666667, 'x[0] <= 36.0\ngini =
0.355\nsamples = 13\nvalue = [3, 10]'),
  Text(0.2524271844660194, 0.2083333333333333, 'gini = 0.5\nsamples =
6\nvalue = [3, 3]'),
  Text(0.27184466019417475, 0.2083333333333333, 'gini = 0.0\nsamples =
7\nvalue = [0, 7]'),
  Text(0.30097087378640774, 0.2916666666666667, 'x[6] <= 13767.5\ngini =
0.375\nsamples = 12\nvalue = [9, 3]'),
  Text(0.2912621359223301, 0.2083333333333333, 'gini = 0.5\nsamples =
6\nvalue = [3, 3]'),
  Text(0.3106796116504854, 0.2083333333333333, 'gini = 0.0\nsamples =
6\nvalue = [6, 0]'),
```

```
Text(0.2621359223300971, 0.4583333333333333, 'gini = 0.245\nsamples = 7\nvalue = [1, 6]'),
Text(0.3446601941747573, 0.7083333333333334, 'x[12] <= 1.5\ngini = 0.279\nsamples = 131\nvalue = [109, 22]'),
Text(0.32038834951456313, 0.625, 'x[20] <= 2.5\ngini = 0.465\nsamples = 19\nvalue = [12, 7]'),
Text(0.3106796116504854, 0.5416666666666666, 'x[8] <= 12.0\ngini = 0.486\nsamples = 12\nvalue = [5, 7]'),
Text(0.30097087378640774, 0.4583333333333333, 'gini = 0.444\nsamples = 6\nvalue = [4, 2]'),
Text(0.32038834951456313, 0.4583333333333333, 'gini = 0.278\nsamples = 6\nvalue = [1, 5]'),
Text(0.3300970873786408, 0.5416666666666666, 'gini = 0.0\nsamples = 7\nvalue = [7, 0]'),
Text(0.36893203883495146, 0.625, 'x[23] <= 2.5\ngini = 0.232\nsamples = 112\nvalue = [97, 15]'),
Text(0.34951456310679613, 0.5416666666666666, 'x[7] <= 94.5\ngini = 0.32\nsamples = 60\nvalue = [48, 12]'),
Text(0.33980582524271846, 0.4583333333333333, 'x[8] <= 26.5\ngini = 0.274\nsamples = 55\nvalue = [46, 9]'),
Text(0.3300970873786408, 0.375, 'x[22] <= 4.5\ngini = 0.215\nsamples = 49\nvalue = [43, 6]'),
Text(0.32038834951456313, 0.2916666666666667, 'gini = 0.48\nsamples = 5\nvalue = [3, 2]'),
Text(0.33980582524271846, 0.2916666666666667, 'x[22] <= 26.5\ngini = 0.165\nsamples = 44\nvalue = [40, 4]'),
Text(0.3300970873786408, 0.20833333333333334, 'gini = 0.0\nsamples = 22\nvalue = [22, 0]'),
Text(0.34951456310679613, 0.20833333333333334, 'x[17] <= 42.5\ngini = 0.298\nsamples = 22\nvalue = [18, 4]'),
Text(0.33980582524271846, 0.125, 'x[22] <= 39.0\ngini = 0.463\nsamples = 11\nvalue = [7, 4]'),
Text(0.3300970873786408, 0.041666666666666664, 'gini = 0.444\nsamples = 6\nvalue = [2, 4]'),
Text(0.34951456310679613, 0.041666666666666664, 'gini = 0.0\nsamples = 5\nvalue = [5, 0]'),
Text(0.3592233009708738, 0.125, 'gini = 0.0\nsamples = 11\nvalue = [11, 0]'),
Text(0.34951456310679613, 0.375, 'gini = 0.5\nsamples = 6\nvalue = [3, 3]'),
Text(0.3592233009708738, 0.4583333333333333, 'gini = 0.48\nsamples = 5\nvalue = [2, 3]'),
Text(0.3883495145631068, 0.5416666666666666, 'x[14] <= 41.5\ngini = 0.109\nsamples = 52\nvalue = [49, 3]'),
Text(0.3786407766990291, 0.4583333333333333, 'gini = 0.0\nsamples = 32\nvalue = [32, 0]'),
Text(0.39805825242718446, 0.4583333333333333, 'x[7] <= 76.0\ngini = 0.255\nsamples = 20\nvalue = [17, 3]'),
Text(0.3883495145631068, 0.375, 'x[22] <= 11.0\ngini = 0.444\nsamples = 9\nvalue = [6, 3]'),
Text(0.3786407766990291, 0.2916666666666667, 'gini = 0.0\nsamples = 4\nvalue = [4, 0]'),
Text(0.39805825242718446, 0.2916666666666667, 'gini = 0.48\nsamples = 5\nvalue = [2, 3]'),
Text(0.4077669902912621, 0.375, 'gini = 0.0\nsamples = 11\nvalue = [11,
```

```

0]'),
  Text(0.5224514563106796, 0.7916666666666666, 'x[8] <= 6.5\ngini =
0.277\nsamples = 247\nvalue = [206, 41]'),
  Text(0.4441747572815534, 0.7083333333333334, 'x[8] <= 2.5\ngini =
0.44\nsamples = 52\nvalue = [35, 17]'),
  Text(0.4174757281553398, 0.625, 'x[10] <= 4.5\ngini = 0.198\nsamples =
18\nvalue = [16, 2]'),
  Text(0.4077669902912621, 0.5416666666666666, 'gini = 0.0\nsamples =
13\nvalue = [13, 0]'),
  Text(0.42718446601941745, 0.5416666666666666, 'gini = 0.48\nsamples =
5\nvalue = [3, 2]'),
  Text(0.470873786407767, 0.625, 'x[9] <= 6.5\ngini = 0.493\nsamples =
34\nvalue = [19, 15]'),
  Text(0.44660194174757284, 0.5416666666666666, 'x[16] <= 4.5\ngini =
0.337\nsamples = 14\nvalue = [11, 3]'),
  Text(0.4368932038834951, 0.4583333333333333, 'gini = 0.5\nsamples = 6\nvalue
= [3, 3]'),
  Text(0.4563106796116505, 0.4583333333333333, 'gini = 0.0\nsamples = 8\nvalue
= [8, 0]'),
  Text(0.49514563106796117, 0.5416666666666666, 'x[10] <= 3.5\ngini =
0.48\nsamples = 20\nvalue = [8, 12]'),
  Text(0.47572815533980584, 0.4583333333333333, 'x[14] <= 81.0\ngini =
0.298\nsamples = 11\nvalue = [2, 9]'),
  Text(0.46601941747572817, 0.375, 'gini = 0.0\nsamples = 6\nvalue = [0, 6]'),
  Text(0.4854368932038835, 0.375, 'gini = 0.48\nsamples = 5\nvalue = [2, 3]'),
  Text(0.5145631067961165, 0.4583333333333333, 'x[15] <= 0.5\ngini =
0.444\nsamples = 9\nvalue = [6, 3]'),
  Text(0.5048543689320388, 0.375, 'gini = 0.375\nsamples = 4\nvalue = [1,
3]'),
  Text(0.5242718446601942, 0.375, 'gini = 0.0\nsamples = 5\nvalue = [5, 0]'),
  Text(0.6007281553398058, 0.7083333333333334, 'x[7] <= 26.5\ngini =
0.216\nsamples = 195\nvalue = [171, 24]'),
  Text(0.5533980582524272, 0.625, 'x[14] <= 89.5\ngini = 0.375\nsamples =
32\nvalue = [24, 8]'),
  Text(0.5436893203883495, 0.5416666666666666, 'x[4] <= 1.5\ngini =
0.252\nsamples = 27\nvalue = [23, 4]'),
  Text(0.5339805825242718, 0.4583333333333333, 'gini = 0.0\nsamples =
14\nvalue = [14, 0]'),
  Text(0.5533980582524272, 0.4583333333333333, 'x[14] <= 71.5\ngini =
0.426\nsamples = 13\nvalue = [9, 4]'),
  Text(0.5436893203883495, 0.375, 'gini = 0.375\nsamples = 4\nvalue = [1,
3]'),
  Text(0.5631067961165048, 0.375, 'x[12] <= 3.5\ngini = 0.198\nsamples =
9\nvalue = [8, 1]'),
  Text(0.5533980582524272, 0.2916666666666667, 'gini = 0.0\nsamples = 5\nvalue
= [5, 0]'),
  Text(0.5728155339805825, 0.2916666666666667, 'gini = 0.375\nsamples =
4\nvalue = [3, 1]'),
  Text(0.5631067961165048, 0.5416666666666666, 'gini = 0.32\nsamples =
5\nvalue = [1, 4]'),
  Text(0.6480582524271845, 0.625, 'x[5] <= 3.5\ngini = 0.177\nsamples =
163\nvalue = [147, 16]'),
  Text(0.6019417475728155, 0.5416666666666666, 'x[21] <= 1.5\ngini =
0.095\nsamples = 100\nvalue = [95, 5]'),

```

```
Text(0.5922330097087378, 0.4583333333333333, 'x[15] <= 0.5\ngini =
0.32\nsamples = 25\nvalue = [20, 5]'),
Text(0.5825242718446602, 0.375, 'gini = 0.0\nsamples = 12\nvalue = [12,
0]'),
Text(0.6019417475728155, 0.375, 'x[18] <= 12.5\ngini = 0.473\nsamples =
13\nvalue = [8, 5]'),
Text(0.5922330097087378, 0.2916666666666667, 'gini = 0.408\nsamples =
7\nvalue = [2, 5]'),
Text(0.6116504854368932, 0.2916666666666667, 'gini = 0.0\nsamples = 6\nvalue
= [6, 0]'),
Text(0.6116504854368932, 0.4583333333333333, 'gini = 0.0\nsamples =
75\nvalue = [75, 0]'),
Text(0.6941747572815534, 0.5416666666666666, 'x[14] <= 76.5\ngini =
0.288\nsamples = 63\nvalue = [52, 11]'),
Text(0.6601941747572816, 0.4583333333333333, 'x[16] <= 4.5\ngini =
0.404\nsamples = 32\nvalue = [23, 9]'),
Text(0.6407766990291263, 0.375, 'x[22] <= 34.0\ngini = 0.133\nsamples =
14\nvalue = [13, 1]'),
Text(0.6310679611650486, 0.2916666666666667, 'gini = 0.0\nsamples =
10\nvalue = [10, 0]'),
Text(0.6504854368932039, 0.2916666666666667, 'gini = 0.375\nsamples =
4\nvalue = [3, 1]'),
Text(0.6796116504854369, 0.375, 'x[16] <= 5.5\ngini = 0.494\nsamples =
18\nvalue = [10, 8]'),
Text(0.6699029126213593, 0.2916666666666667, 'gini = 0.0\nsamples = 5\nvalue
= [0, 5]'),
Text(0.6893203883495146, 0.2916666666666667, 'x[22] <= 18.0\ngini =
0.355\nsamples = 13\nvalue = [10, 3]'),
Text(0.6796116504854369, 0.2083333333333333, 'gini = 0.48\nsamples =
5\nvalue = [2, 3]'),
Text(0.6990291262135923, 0.2083333333333333, 'gini = 0.0\nsamples =
8\nvalue = [8, 0]'),
Text(0.7281553398058253, 0.4583333333333333, 'x[6] <= 6698.5\ngini =
0.121\nsamples = 31\nvalue = [29, 2]'),
Text(0.7184466019417476, 0.375, 'x[6] <= 4182.5\ngini = 0.375\nsamples =
8\nvalue = [6, 2]'),
Text(0.7087378640776699, 0.2916666666666667, 'gini = 0.0\nsamples = 4\nvalue
= [4, 0]'),
Text(0.7281553398058253, 0.2916666666666667, 'gini = 0.5\nsamples = 4\nvalue
= [2, 2]'),
Text(0.7378640776699029, 0.375, 'gini = 0.0\nsamples = 23\nvalue = [23,
0]'),
Text(0.9235436893203883, 0.875, 'x[18] <= 18.5\ngini = 0.232\nsamples =
247\nvalue = [214, 33]'),
Text(0.7864077669902912, 0.9166666666666667, ' False'),
Text(0.866504854368932, 0.7916666666666666, 'x[22] <= 21.5\ngini =
0.207\nsamples = 230\nvalue = [203, 27]'),
Text(0.8155339805825242, 0.7083333333333333, 'x[12] <= 3.5\ngini =
0.299\nsamples = 104\nvalue = [85, 19]'),
Text(0.7766990291262136, 0.625, 'x[16] <= 7.5\ngini = 0.192\nsamples =
65\nvalue = [58, 7]'),
Text(0.7572815533980582, 0.5416666666666666, 'x[22] <= 17.5\ngini =
0.078\nsamples = 49\nvalue = [47.0, 2.0]'),
Text(0.7475728155339806, 0.4583333333333333, 'gini = 0.0\nsamples =
```

```


38\nvalue = [38, 0]'),
  Text(0.7669902912621359, 0.4583333333333333, 'x[3] <= 1.5\ngini =
0.298\nsamples = 11\nvalue = [9, 2]'),
  Text(0.7572815533980582, 0.375, 'gini = 0.5\nsamples = 4\nvalue = [2, 2]'),
  Text(0.7766990291262136, 0.375, 'gini = 0.0\nsamples = 7\nvalue = [7, 0]'),
  Text(0.7961165048543689, 0.5416666666666666, 'x[18] <= 5.5\ngini =
0.43\nsamples = 16\nvalue = [11, 5]'),
  Text(0.7864077669902912, 0.4583333333333333, 'gini = 0.32\nsamples =
5\nvalue = [1, 4]'),
  Text(0.8058252427184466, 0.4583333333333333, 'x[6] <= 7828.5\ngini =
0.165\nsamples = 11\nvalue = [10, 1]'),
  Text(0.7961165048543689, 0.375, 'gini = 0.375\nsamples = 4\nvalue = [3,
1]'),
  Text(0.8155339805825242, 0.375, 'gini = 0.0\nsamples = 7\nvalue = [7, 0]'),
  Text(0.8543689320388349, 0.625, 'x[9] <= 11.5\ngini = 0.426\nsamples =
39\nvalue = [27, 12]'),
  Text(0.8446601941747572, 0.5416666666666666, 'x[7] <= 32.0\ngini =
0.469\nsamples = 32\nvalue = [20, 12]'),
  Text(0.8349514563106796, 0.4583333333333333, 'gini = 0.408\nsamples =
7\nvalue = [2, 5]'),
  Text(0.8543689320388349, 0.4583333333333333, 'x[11] <= 1.5\ngini =
0.403\nsamples = 25\nvalue = [18.0, 7.0]'),
  Text(0.8349514563106796, 0.375, 'x[6] <= 13835.0\ngini = 0.494\nsamples =
9\nvalue = [4, 5]'),
  Text(0.8252427184466019, 0.2916666666666667, 'gini = 0.32\nsamples =
5\nvalue = [1, 4]'),
  Text(0.8446601941747572, 0.2916666666666667, 'gini = 0.375\nsamples =
4\nvalue = [3, 1]'),
  Text(0.8737864077669902, 0.375, 'x[19] <= 3.5\ngini = 0.219\nsamples =
16\nvalue = [14, 2]'),
  Text(0.8640776699029126, 0.2916666666666667, 'gini = 0.0\nsamples =
11\nvalue = [11, 0]'),
  Text(0.883495145631068, 0.2916666666666667, 'gini = 0.48\nsamples = 5\nvalue
= [3, 2]'),
  Text(0.8640776699029126, 0.5416666666666666, 'gini = 0.0\nsamples = 7\nvalue
= [7, 0]'),
  Text(0.9174757281553398, 0.7083333333333334, 'x[17] <= 31.5\ngini =
0.119\nsamples = 126\nvalue = [118, 8]'),
  Text(0.8932038834951457, 0.625, 'x[18] <= 8.5\ngini = 0.469\nsamples =
8\nvalue = [5, 3]'),
  Text(0.883495145631068, 0.5416666666666666, 'gini = 0.375\nsamples =
4\nvalue = [3, 1]'),
  Text(0.9029126213592233, 0.5416666666666666, 'gini = 0.5\nsamples = 4\nvalue
= [2, 2]'),
  Text(0.941747572815534, 0.625, 'x[10] <= 8.5\ngini = 0.081\nsamples =
118\nvalue = [113, 5]'),
  Text(0.9223300970873787, 0.5416666666666666, 'x[6] <= 18745.5\ngini =
0.025\nsamples = 80\nvalue = [79, 1]'),
  Text(0.912621359223301, 0.4583333333333333, 'gini = 0.0\nsamples = 73\nvalue
= [73, 0]'),
  Text(0.9320388349514563, 0.4583333333333333, 'gini = 0.245\nsamples =
7\nvalue = [6, 1]'),
  Text(0.9611650485436893, 0.5416666666666666, 'x[17] <= 49.0\ngini =
0.188\nsamples = 38\nvalue = [34, 4]'),

```

```

Text(0.9514563106796117, 0.4583333333333333, 'gini = 0.0\nsamples =
22\nvalue = [22, 0]'),
Text(0.970873786407767, 0.4583333333333333, 'x[12] <= 3.5\ngini =
0.375\nsamples = 16\nvalue = [12, 4]'),
Text(0.9611650485436893, 0.375, 'x[4] <= 1.5\ngini = 0.153\nsamples =
12\nvalue = [11, 1]'),
Text(0.9514563106796117, 0.2916666666666667, 'gini = 0.375\nsamples =
4\nvalue = [3, 1]'),
Text(0.970873786407767, 0.2916666666666667, 'gini = 0.0\nsamples = 8\nvalue
= [8, 0]'),
Text(0.9805825242718447, 0.375, 'gini = 0.375\nsamples = 4\nvalue = [1,
3]'),
Text(0.9805825242718447, 0.7916666666666666, 'x[7] <= 80.5\ngini =
0.457\nsamples = 17\nvalue = [11, 6]'),
Text(0.970873786407767, 0.7083333333333334, 'x[13] <= 3.5\ngini =
0.26\nsamples = 13\nvalue = [11, 2]'),
Text(0.9611650485436893, 0.625, 'gini = 0.0\nsamples = 9\nvalue = [9, 0]'),
Text(0.9805825242718447, 0.625, 'gini = 0.5\nsamples = 4\nvalue = [2, 2]'),
Text(0.9902912621359223, 0.7083333333333334, 'gini = 0.0\nsamples = 4\nvalue
= [0, 4]')

```



```

# Obtener el accuracy_score
y_pred_c = tree_clf_c.predict(X_test)
accuracy_c = accuracy_score(y_test, y_pred_c)
print("El accuracy score que se tiene es %s" % accuracy_c)

```

El accuracy score que se tiene es 0.71


```

# Obtener la matriz de confusión
cm = confusion_matrix(y_test, y_pred_c)
sns.heatmap(cm, annot=True, fmt="d")

```

```
plt.title("Árbol de Decisiones - min_samples_leaf=4")
plt.xlabel("Predicho")
plt.ylabel("Real")
```

```
Text(50.72222222222214, 0.5, 'Real')
```



```
# Obtener la precisión, recall y f1
```

```
precision_c = precision_score(y_test, y_pred_c, average="macro")
```

```
recall_c = recall_score(y_test, y_pred_c, average="macro")
```

```
f1_c = f1_score(y_test, y_pred_c, average="macro")
```

```
metricas = pd.DataFrame({"Métricas": ["Precisión", "Recall", "F1"],
                          "Valores": [precision_c, recall_c, f1_c]})
```

```
metricas
```

	Métricas	Valores
0	Precisión	0.542857
1	Recall	0.536661
2	F1	0.538143

Comparación de los modelos de Árbol de Decisiones

```
comparacion_tree = pd.DataFrame({
    'Modelo': ['Árbol - Defecto', 'Árbol - max_depth=5', 'Árbol - min_samples_leaf=4'],
```

```

'Accuracy': [accuracy_a, accuracy_b, accuracy_c],
'Precisión': [precision_a, precision_b, precision_c],
'Recall': [recall_a, recall_b, recall_c],
'F1': [f1_a, f1_b, f1_c]
})
comparacion_tree

```

	Modelo	Accuracy	Precisión	Recall	F1
0	Árbol - Defecto	0.680	0.564890	0.576655	0.567509
1	Árbol - max_depth=5	0.765	0.475086	0.495704	0.453457
2	Árbol - min_samples_leaf=4	0.710	0.542857	0.536661	0.538143

7. Comparación Global de todos los Algoritmos

A continuación se presenta un resumen consolidado de todos los modelos entrenados, permitiendo identificar el mejor algoritmo para la detección temprana de abandono laboral.

```

comparacion_global = pd.DataFrame({
    'Algoritmo': [
        'Regresión Logística (defecto)',
        'Regresión Logística (L1)',
        'Regresión Logística (C=0.5)',
        'KNN (k=5)',
        'KNN (k=3)',
        'KNN (k=11)',
        'Árbol de Decisión (defecto)',
        'Árbol de Decisión (max_depth=5)',
        'Árbol de Decisión (min_samples_leaf=4)'
    ],
    'Accuracy': [
        accuracy_lr_a, accuracy_lr_b, accuracy_lr_c,
        accuracy_knn_a, accuracy_knn_b, accuracy_knn_c,
        accuracy_a, accuracy_b, accuracy_c
    ],
    'F1 (macro)': [
        f1_lr_a, f1_lr_b, f1_lr_c,
        f1_knn_a, f1_knn_b, f1_knn_c,
        f1_a, f1_b, f1_c
    ]
})

comparacion_global.sort_values('F1 (macro)',
                               ascending=False).reset_index(drop=True)


```

	Algoritmo	Accuracy	F1 (macro)
0	Árbol de Decisión (defecto)	0.680	0.567509

	Algoritmo	Accuracy	F1 (macro)
1	Árbol de Decisión (min_samples_leaf=4)	0.710	0.538143
2	KNN (k=3)	0.740	0.475701
3	KNN (k=5)	0.770	0.455492
4	Árbol de Decisión (max_depth=5)	0.765	0.453457
5	Regresión Logística (defecto)	0.785	0.439776
6	Regresión Logística (C=0.5)	0.785	0.439776
7	Regresión Logística (L1)	0.785	0.439776
8	KNN (k=11)	0.780	0.438202

Gráfico comparativo de Accuracy por algoritmo

```
plt.figure(figsize=(12, 5))
colores = ['#2196F3'] * 3 + ['#4CAF50'] * 3 + ['#FF9800'] * 3
barras = plt.bar(range(9), comparacion_global['Accuracy'], color=colores)
for barra in barras:
    altura = barra.get_height()
    plt.text(barra.get_x() + barra.get_width() / 2, altura, f'{altura:.3f}',
             ha='center', va='bottom')
plt.xticks(range(9), comparacion_global['Algoritmo'], rotation=45,
           ha='right', fontsize=9)
plt.ylabel('Accuracy')
plt.title('Comparación de Accuracy por Algoritmo de Clasificación')
plt.ylim(0, 1)
plt.tight_layout()
plt.show()
```



#Comparación global de todas las métricas

```
comparacion_metricas_global = pd.DataFrame({
    'Modelo': [
        'RL - Defecto',
        'RL - L1 (Lasso)',
        'RL - C=0.5, iter=500',
        'KNN - k=5',
        'KNN - k=3',
    ]
})
```


```

        'KNN - k=11',
        'Árbol - Defecto',
        'Árbol - max_depth=5',
        'Árbol - min_samples_leaf=4'
    ],
    'Accuracy': [
        accuracy_lr_a, accuracy_lr_b, accuracy_lr_c,
        accuracy_knn_a, accuracy_knn_b, accuracy_knn_c,
        accuracy_a, accuracy_b, accuracy_c
    ],
    'Precisión': [
        precision_lr_a, precision_lr_b, precision_lr_c,
        precision_knn_a, precision_knn_b, precision_knn_c,
        precision_a, precision_b, precision_c
    ],
    'Recall': [
        recall_lr_a, recall_lr_b, recall_lr_c,
        recall_knn_a, recall_knn_b, recall_knn_c,
        recall_a, recall_b, recall_c
    ],
    'F1': [
        f1_lr_a, f1_lr_b, f1_lr_c,
        f1_knn_a, f1_knn_b, f1_knn_c,
        f1_a, f1_b, f1_c
    ]
})

comparacion_metricas_larga = comparacion_metricas_global.melt(
    id_vars='Modelo',
    var_name='Métrica',
    value_name='Valor'
)

plt.figure(figsize=(16, 6))
ax = sns.barplot(x='Modelo', y='Valor', hue='Métrica',
    data=comparacion_metricas_larga)
for p in ax.patches:
    valor = p.get_height()
    ax.annotate(f'{valor:.3f}',
        (p.get_x() + p.get_width() / 2, valor),
        ha='center', va='bottom', fontsize=8, rotation=45)
plt.title('Comparación global de métricas por modelo')
plt.xlabel('Modelo')
plt.ylabel('Valor')
plt.ylim(0, 1)
plt.xticks(rotation=45, ha='right')
plt.legend(title='Métrica')
plt.tight_layout()
plt.show()

```



8. Conclusiones

Con base en los resultados obtenidos a partir de la aplicación de los tres algoritmos de clasificación (Regresión Logística, KNN y Árbol de Decisiones) sobre el conjunto de datos de abandono laboral, se pueden extraer las siguientes conclusiones:

1. **Regresión Logística:** Ofrece un rendimiento sólido y estable entre sus diferentes configuraciones. La regularización permite controlar el sobreajuste y mejorar la generalización del modelo.
2. **KNN:** Es sensible al valor de k. Valores muy pequeños de k pueden provocar sobreajuste, mientras que valores muy grandes pueden suavizar en exceso las fronteras de decisión. Es importante normalizar los datos antes de aplicar este algoritmo.
3. **Árbol de Decisiones:** Sin restricciones puede sobreajustarse completamente a los datos de entrenamiento. La limitación de la profundidad máxima (max_depth) o del número mínimo de muestras por hoja (min_samples_leaf) ayuda a construir modelos más generalizables.

La selección del mejor modelo debe considerar no solo el accuracy, sino también métricas como F1, precisión y recall, especialmente cuando existe desbalance de clases en la variable objetivo Attrition.